

Victor Isuo

Applied LLM Systems Engineer | RAG · Multi-Agent Systems · LLM Evaluation · AWS

victorisuo@proton.me · victor-isuo.github.io · linkedin.com/in/victor-isuo-a02b65171 · github.com/victor-isuo

Nigeria (UTC+1) · Open to remote

SUMMARY

Applied LLM systems engineer. Builds RAG and multi-agent systems, deploys them on AWS, and gates releases on automated evaluation. Four systems shipped: a multi-agent diagnosis copilot on AWS (0.90 mean on a 30-case eval, 3.1 s average latency), an evaluation platform that blocks deploys on 5% regressions, an autonomous intelligence pipeline, and a contract-risk analyzer.

TECHNICAL SKILLS

LLM & Agents: LangGraph · LangChain · CrewAI · MCP · tool calling · multi-agent orchestration · Groq (Llama 4) · Gemini 2.5 Flash

Retrieval: Hybrid search (BM25 + dense) · ChromaDB · Cohere reranking · Sentence Transformers · chunking strategy

Evaluation: LLM-as-judge · custom eval suites · regression gating · LangSmith tracing

Cloud & DevOps: AWS (ECS Fargate, ALB, ECR, S3, Secrets Manager, IAM) · Docker · GitHub Actions (OIDC) · CI/CD

Backend: Python · FastAPI · Pydantic · asyncio · REST APIs · SQLite · MQTT · APScheduler

PROJECTS

Industrial AI Copilot | *Multi-agent fault diagnosis platform, deployed on AWS*

github.com/victor-isuo/industrial-ai-copilot

- Designed hybrid retrieval over 27 industrial documents (5,091 chunks): BM25 and dense embeddings fused 0.5/0.5, top 8 candidates reranked by Cohere, 512/100 chunk size/overlap tuned for exact standards and part codes (ISO, NFPA). Every answer cites document and page.
- Implemented a 9-tool LangGraph agent (spec checking, live MQTT telemetry, Gemini vision, unit conversion, MCP client) and a supervisor that routes cross-domain queries to 4 specialist agents with isolated context windows; a report agent merges findings (full-plant audit in ~30 s).
- Exposed the tool set as an MCP server (connectable from Claude Desktop) and traced every tool call, latency, and token count in LangSmith.
- Evaluated on 30 hand-written cases (spec checks, unit conversion, retrieval, edge cases) authored before any optimization. Scoring: tool selection 40%, keyword match 40%, severity 20%; pass at 0.70. Result: 0.90 mean, 3.1 s average latency; each failure traced to a root cause and fix.
- Migrated from Hugging Face Spaces to AWS: ECS Fargate behind an ALB, images in ECR, ChromaDB restored from S3, keys in Secrets Manager, GitHub Actions deploying via OIDC. Fixed a health-check race with a readiness gate (503 until loaded).

AgentEval | *LLM evaluation and observability platform*

github.com/victor-isuo/agenteval

- Scored RAG, single-agent, and multi-agent systems with LLM-as-judge on weighted axes (RAG: faithfulness 40%, citation accuracy 35%, relevance 25%) across 25 cases; pass at 0.70.
- Gated deploys in GitHub Actions: pytest, then a 5-case smoke eval compared against the previous run of the same mode; a drop over 5% blocks release. Results stored in SQLite; dashboard shows score trends and failure root causes.

Busiv | *Autonomous business intelligence pipeline*

github.com/victor-isuo/busiv

- Automated ingestion (APScheduler): 9 RSS feeds every 6 h, relevance-scored against a watchlist (10 companies, 4 regulators), deduplicated by SHA256, indexed in ChromaDB. A daily 3-node LangGraph step synthesizes cited briefings by signal type (regulatory, product, market, hiring) with priority alerts, sent by dashboard and email.
- Domain set in one config file; demonstrated on Nigerian fintech with CBN and SEC monitoring.

LexAI | *AI contract risk analysis*

github.com/victor-isuo/lexai

- Parsed contract PDFs into typed clauses (boundary detection, 14 regex risk signals) and passed the highest-risk first through a 4-agent CrewAI chain (extract, research, risk-score, report) grounded in a curated Common Law/UNIDROIT knowledge base.
- Scored every run with an independent LLM judge (clause detection 40%, legal grounding 35%, recommendation specificity 25%; pass at 0.70); fallback LLM chain and TPM-aware retries handle rate limits.

EDUCATION

IBM RAG & Agentic AI Professional Certificate | IBM / Coursera

Feb 2026

Microsoft AI & ML Engineering Professional Certificate | Microsoft / Coursera

Oct 2025

B.Eng. Mechanical Engineering | University of Uyo, Nigeria

Graduated March 2025